

# homework-4

```
library(bis557)
```

## Question 1

In Python, implement a numerically-stable ridge regression that takes into account colinear (or nearly colinear) regression variables. Show that it works by comparing it to the output of your R implementation.

The R implementation:

```
data(iris)
iris$target = ifelse(iris$Species=="setosa", 0, ifelse(iris$Species=="versicolor", 1, 2))
X = iris[,1:3]
y = iris$target
ridge(X, y, 0.01)
#> $coefficients
#>           [,1]
#> [1,] -0.23152809
#> [2,]  0.09850707
#> [3,]  0.54588106
#>
#> $lambda
#>           [,1] [,2] [,3]
#> [1,] 0.01 0.00 0.00
#> [2,] 0.00 0.01 0.00
#> [3,] 0.00 0.00 0.01
```

The python output (notebook) shows that the coefficients are -0.23178751, 0.09881172 and 0.54603559, which are exactly the same as R output.

## Question 2

Create an “out-of-core” implementation of the linear model that reads in contiguous rows of a data frame from a file, updates the model. You may read the data from R and send it to your Python functions for fitting. (please see python notebook)

## Question 3

Implement your own LASSO regression function in Python. Show that the results are the same as the function implemented in the `casl` package.

```
casl_util_soft_thresh <-
function(a, b)
{
  a[abs(a) <= b] <- 0
  a[a > 0] <- a[a > 0] - b
  a[a < 0] <- a[a < 0] + b
  a
}
```

```

}

casl_lenet_update_beta <-
function(X, y, lambda, alpha, b, W)
{
  WX <- W * X
  WX2 <- W * X^2
  Xb <- X %*% b

  for (i in seq_along(b))
  {
    Xb <- Xb - X[, i] * b[i]
    b[i] <- casl_util_soft_thresh(sum(WX[,i, drop=FALSE] *
                                   (y - Xb)),
                                   lambda*alpha)
    b[i] <- b[i] / (sum(WX2[, i]) + lambda * (1 - alpha))
    Xb <- Xb + X[, i] * b[i]
  }
  b
}

# the casl lasso function
casl_lenet <-
function(X, y, lambda, alpha = 1, b=matrix(0, nrow=ncol(X), ncol=1),
          tol = 1e-5, maxit=50L, W=rep(1, length(y))/length(y))
{
  for (j in seq_along(lambda))
  {
    if (j > 1)
    {
      b[,j] <- b[, j-1, drop = FALSE]
    }

    # Update the slope coefficients until they converge.
    for (i in seq(1, maxit))
    {
      b_old <- b[, j]
      b[, j] <- casl_lenet_update_beta(X, y, lambda[j], alpha,
                                       b[, j], W)

      if (all(abs(b[, j] - b_old) < tol)) {
        break
      }
    }
    if (i == maxit)
    {
      warning("Function lenet did not converge.")
    }
  }
  b
}

```

```

#check consistency
diabetes = read.table("https://web.stanford.edu/~hastie/Papers/LARS/diabetes.data",
  header = T)
X = as.matrix(diabetes[, 1:10])
y = diabetes$Y

casl_lenet(X, y, lambda=1)
#> Warning in casl_lenet(X, y, lambda = 1): Function lenet did not converge.
#>      [,1]
#> [1,]  0.3279708
#> [2,] -19.6516085
#> [3,]  7.1969435
#> [4,]  0.3035790
#> [5,] -0.2384569
#> [6,] -0.3613295
#> [7,] -0.9590866
#> [8,] 14.5782324
#> [9,] 10.1761675
#> [10,] -0.2634052

```

The results are the same as the LASSO function in python.

#### Question 4

**Topic:** optimize the  $K$  parameter in  $k$ -nearest neighbors(KNN) algorithm via cross-validation.

**Benchmark comparison:** ROC (Receiver Operating Characteristic) Curve to evaluate the diagnostic ability of binary classifiers.

**Application:** predict the outcome of diabetes using the Pima Indians Diabetes Database.